

Introduction to ANN and Perceptron

CS1090B Data Science II – Spring 2025
Pavlos Protopapas, Natesh Pillai, and Chris Gumb

Outline

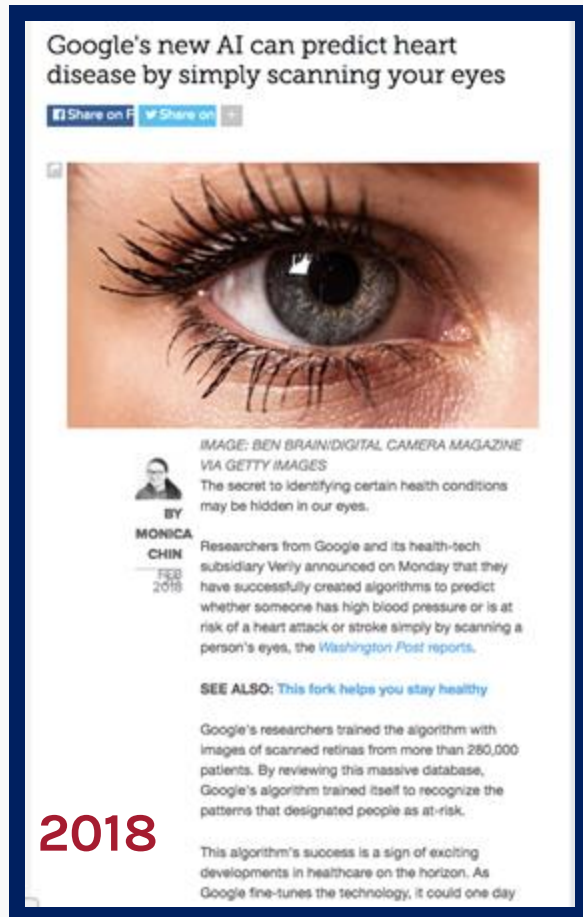
- Introduction
- Review of basic concepts
- Linear to Logistic Regression
- Logistic Regression to ANN
- Perceptron - Single neuron network
- Multi-Layer Perceptron (MLP)

Outline

- **Introduction**
- Review of basic concepts
- Linear to Logistic Regression
- Logistic Regression to ANN
- Perceptron - Single neuron network
- Multi-Layer Perceptron (MLP)

Historical Trends

Disease prediction

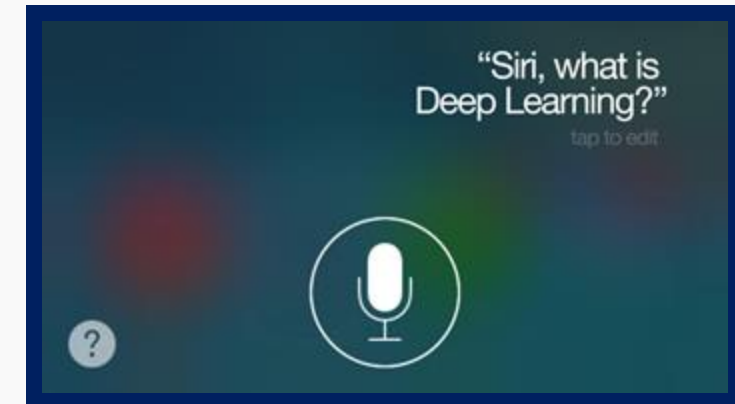


2018

Game strategy




Natural Language Processing



2011

Recent developments

Protein folding



AlphaFold, a DeepMind AI, revolutionized biochemistry by solving the long-standing protein folding problem.

2020

Autonomous cars



AI Detecting objects to assist with autonomous driving.


Recent developments

Image Reconstruction from Sparse Frequency Measurements



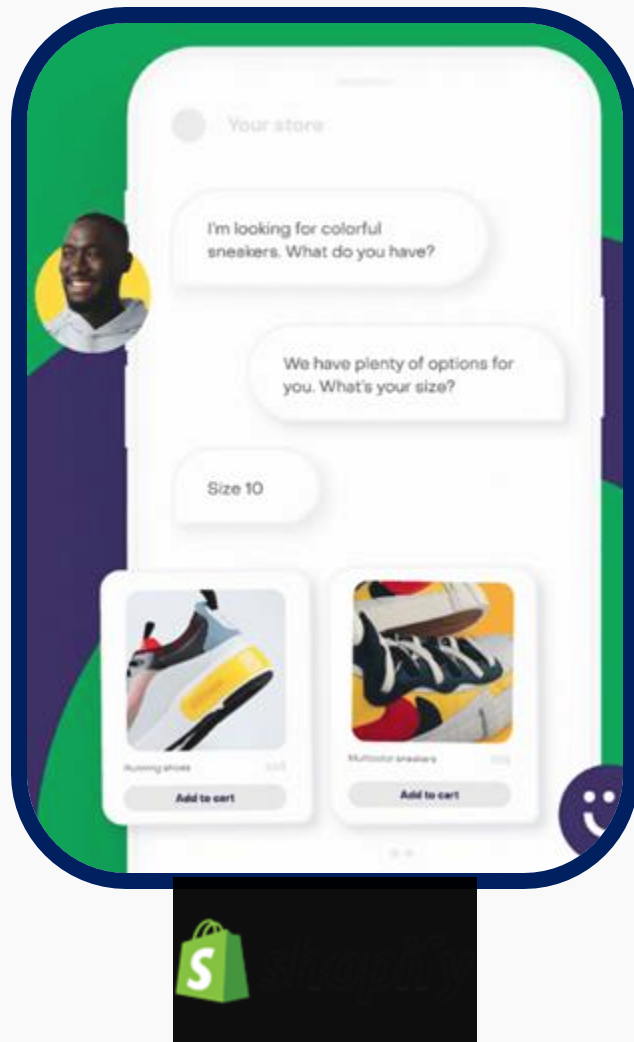
Katie Bouman's CHIRP produces the first-ever image of a black hole.

Text to Image Generation

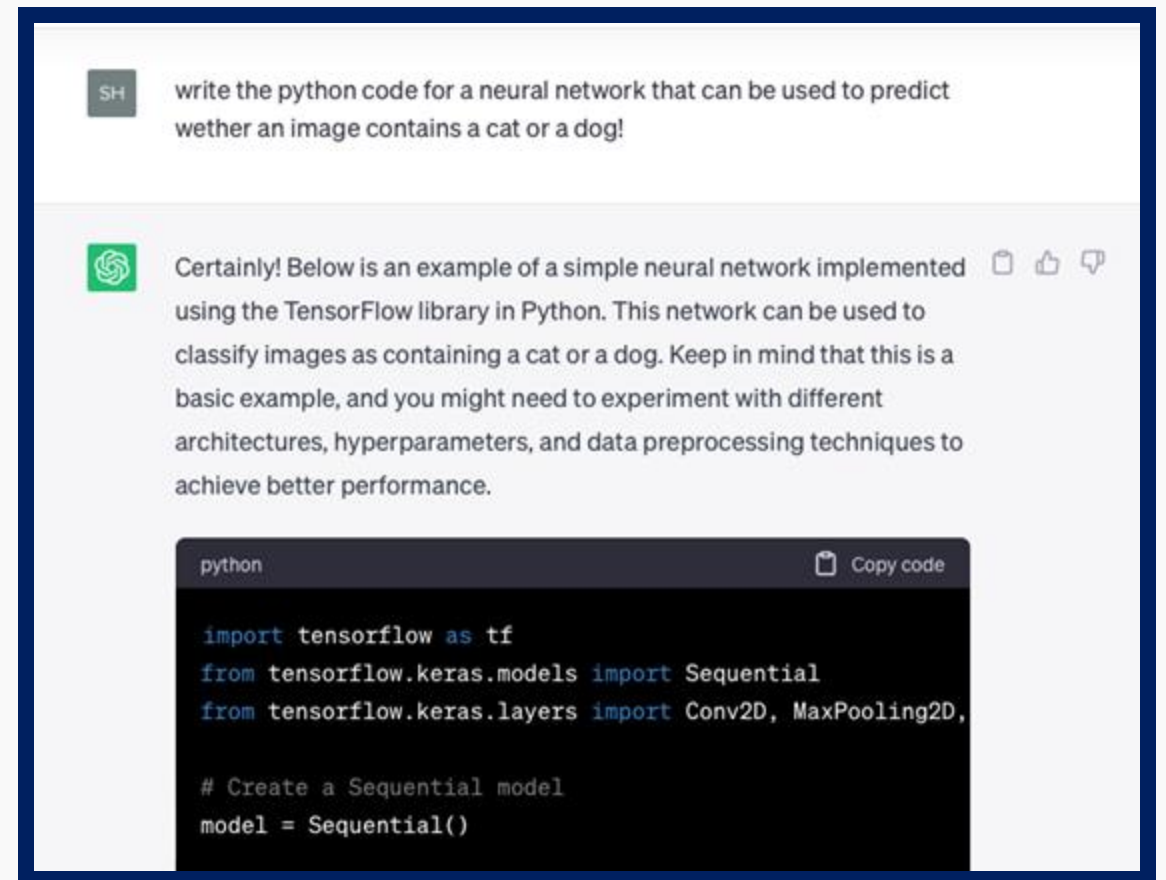


Recent developments

Personalized Customer Assistance




Computer Code Generation



Recent developments

Disease Prediction



Google, 2023

Complex Object Detection



YOLOv5, 2024

The potential challenges in Data Science

Gender Bias



Some DS models for evaluating job applications show bias in favor of male candidate

Racial Bias




Risk models used in US courts have shown to be biased against non-white defendants


Outline

- Introduction
- **Review of basic concepts**
- Linear to Logistic Regression
- Logistic Regression to ANN
- Perceptron - Single neuron network
- Multi-Layer Perceptron (MLP)

Supervised v/s Unsupervised Machine Learning

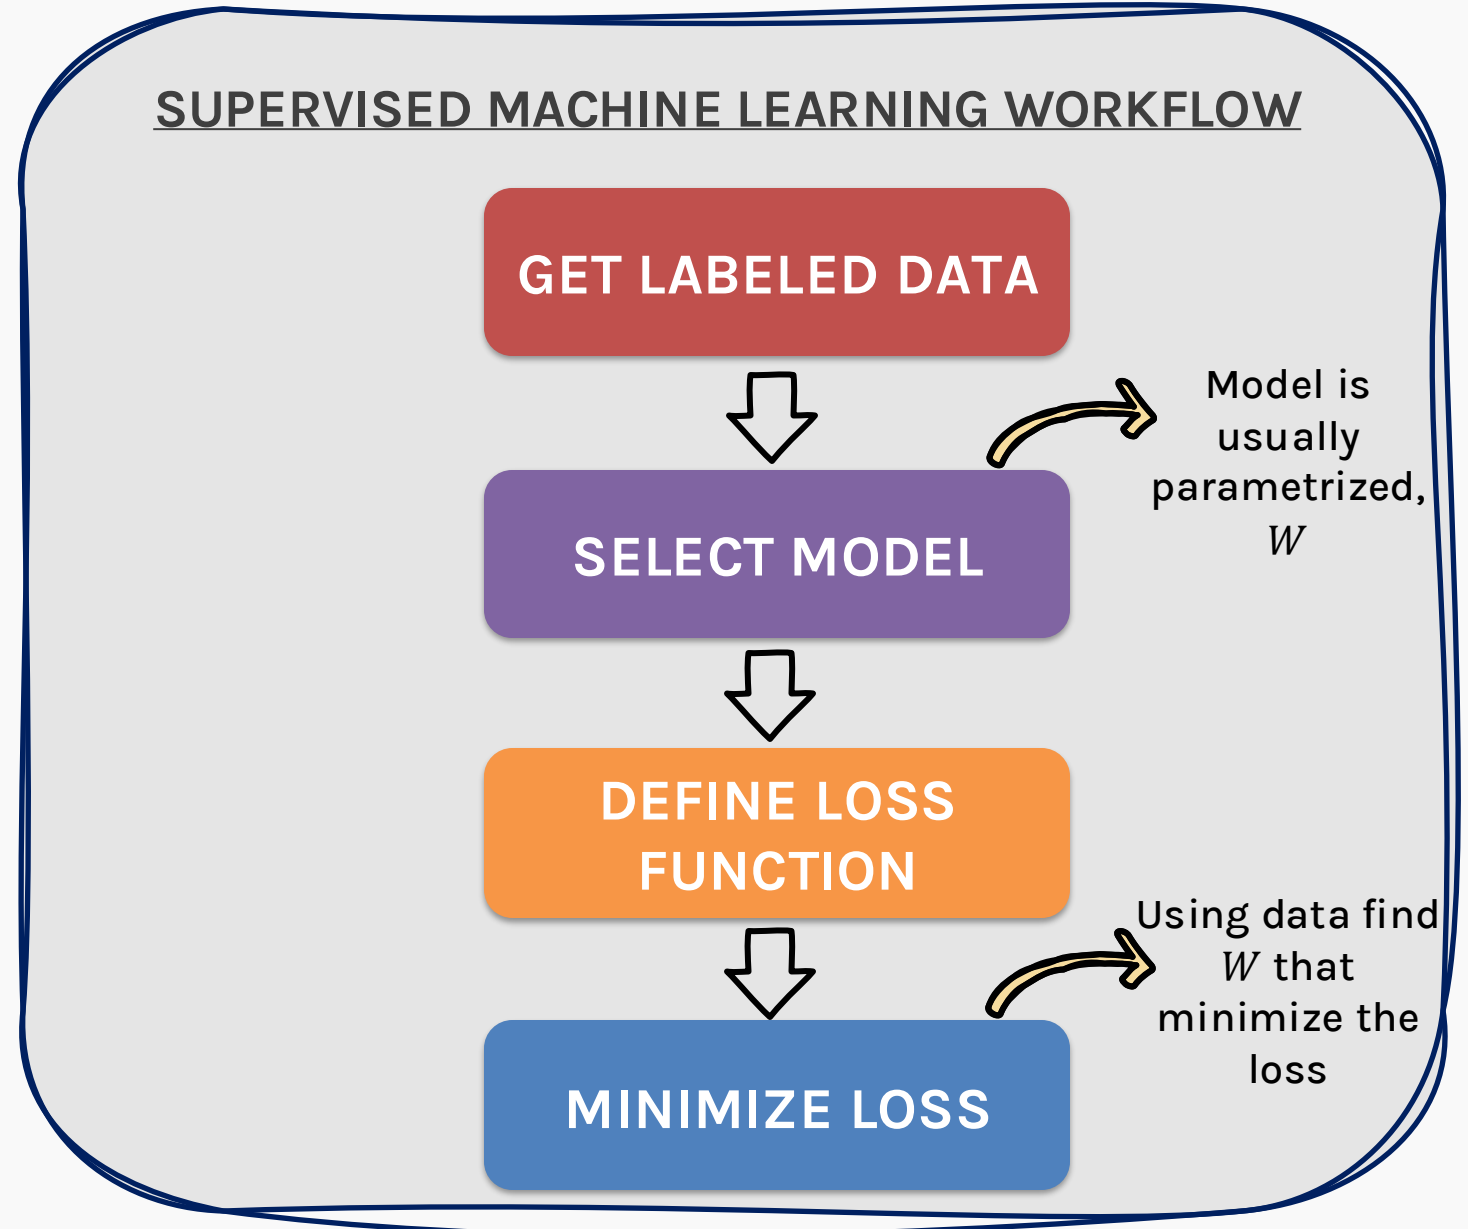


Supervised Learning: Learns with “labeled” data



Unsupervised Learning: Learns by clustering or association

Building blocks of supervised machine learning

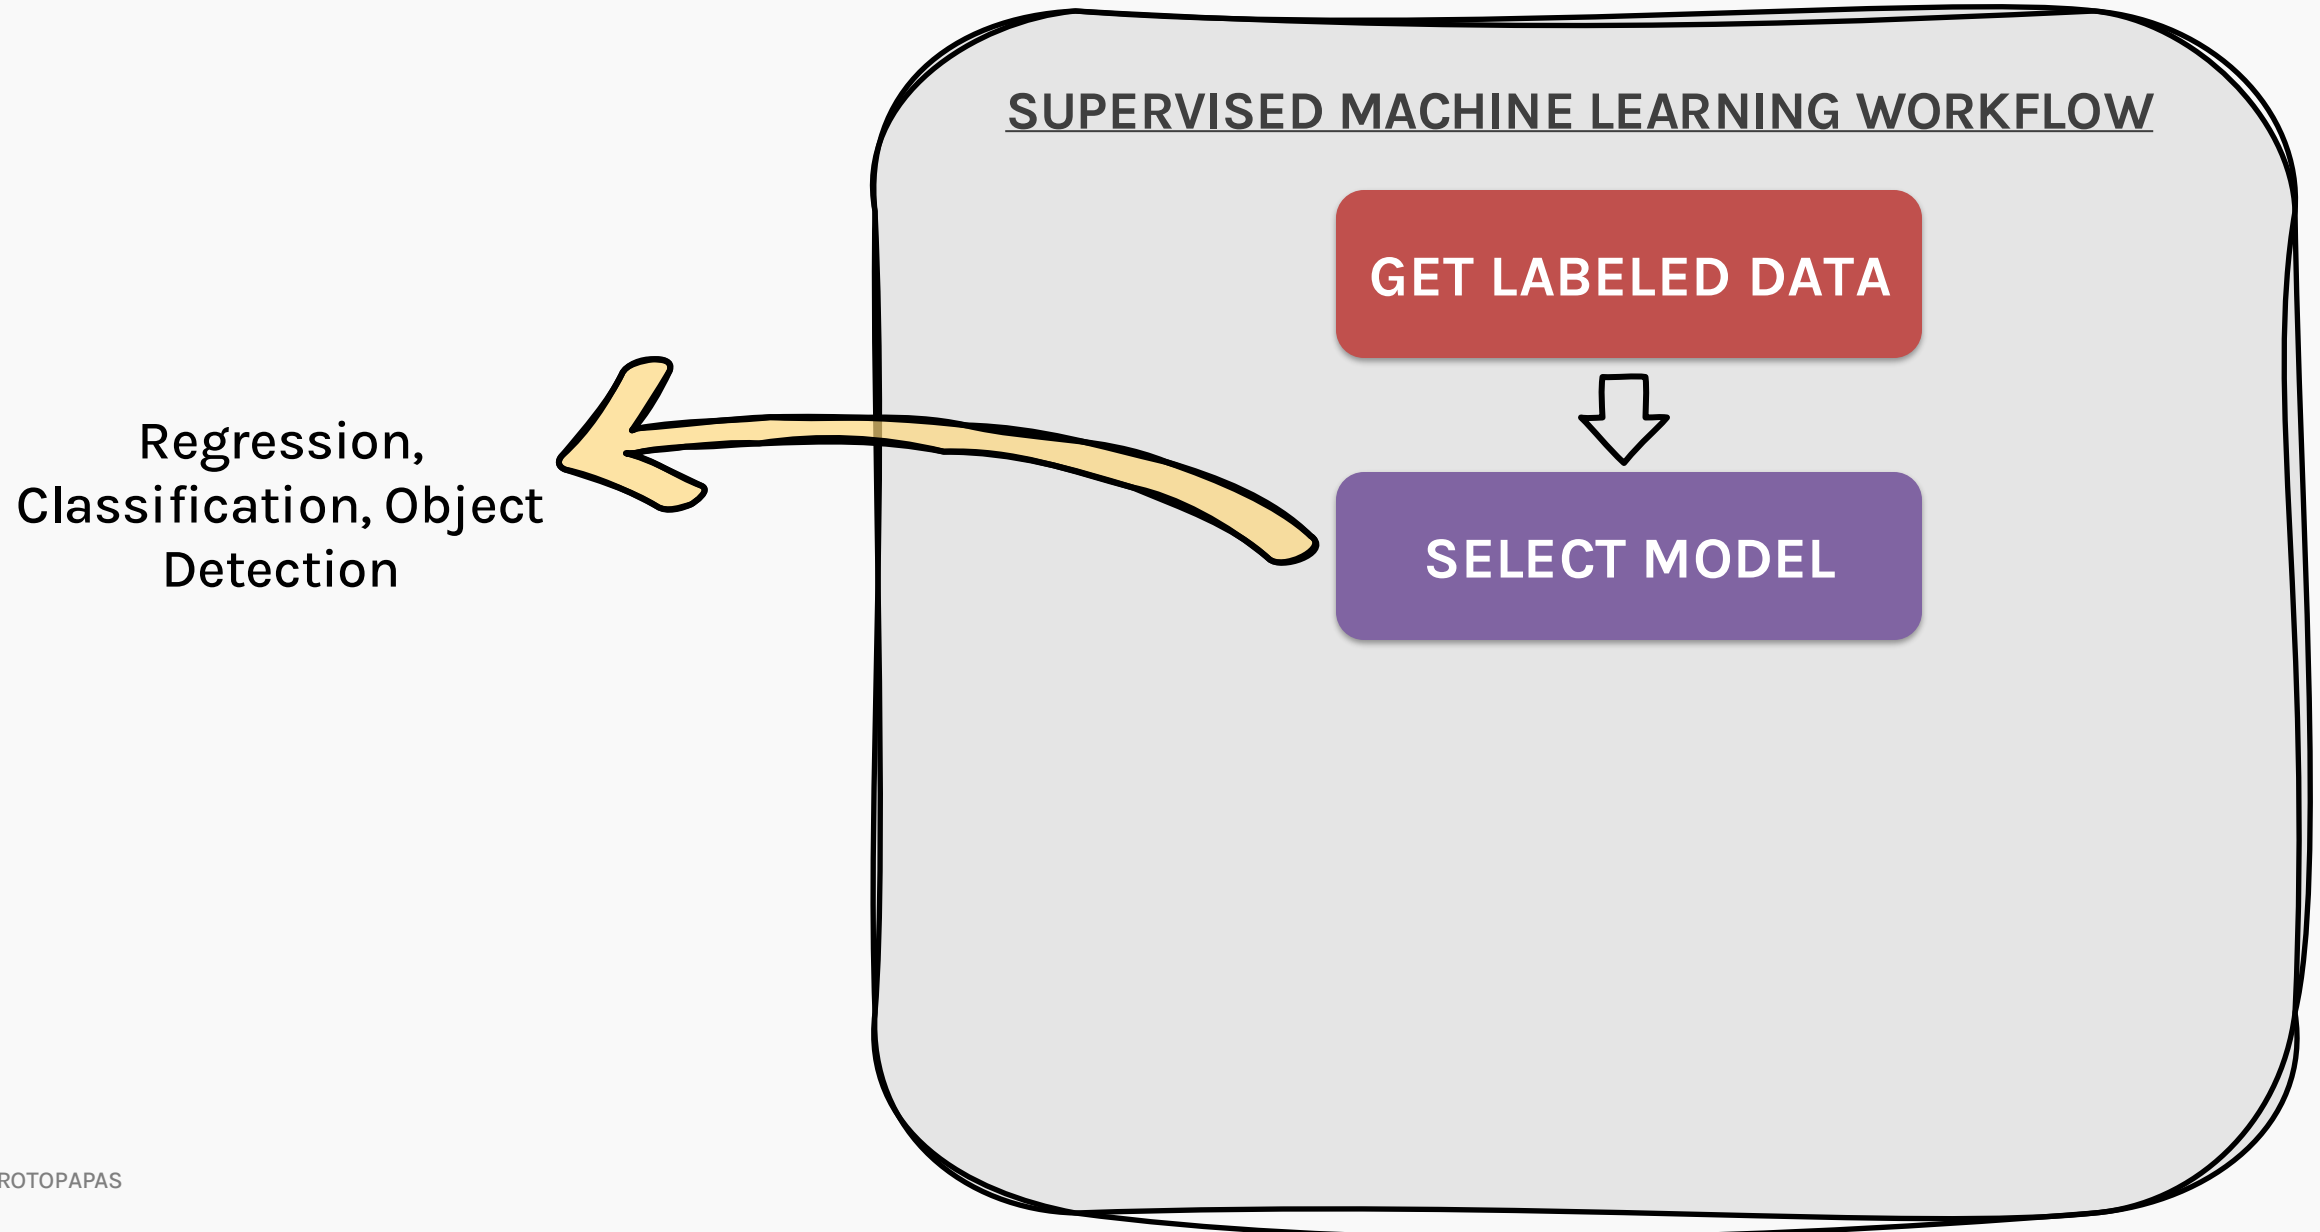


Building blocks of **supervised** machine learning

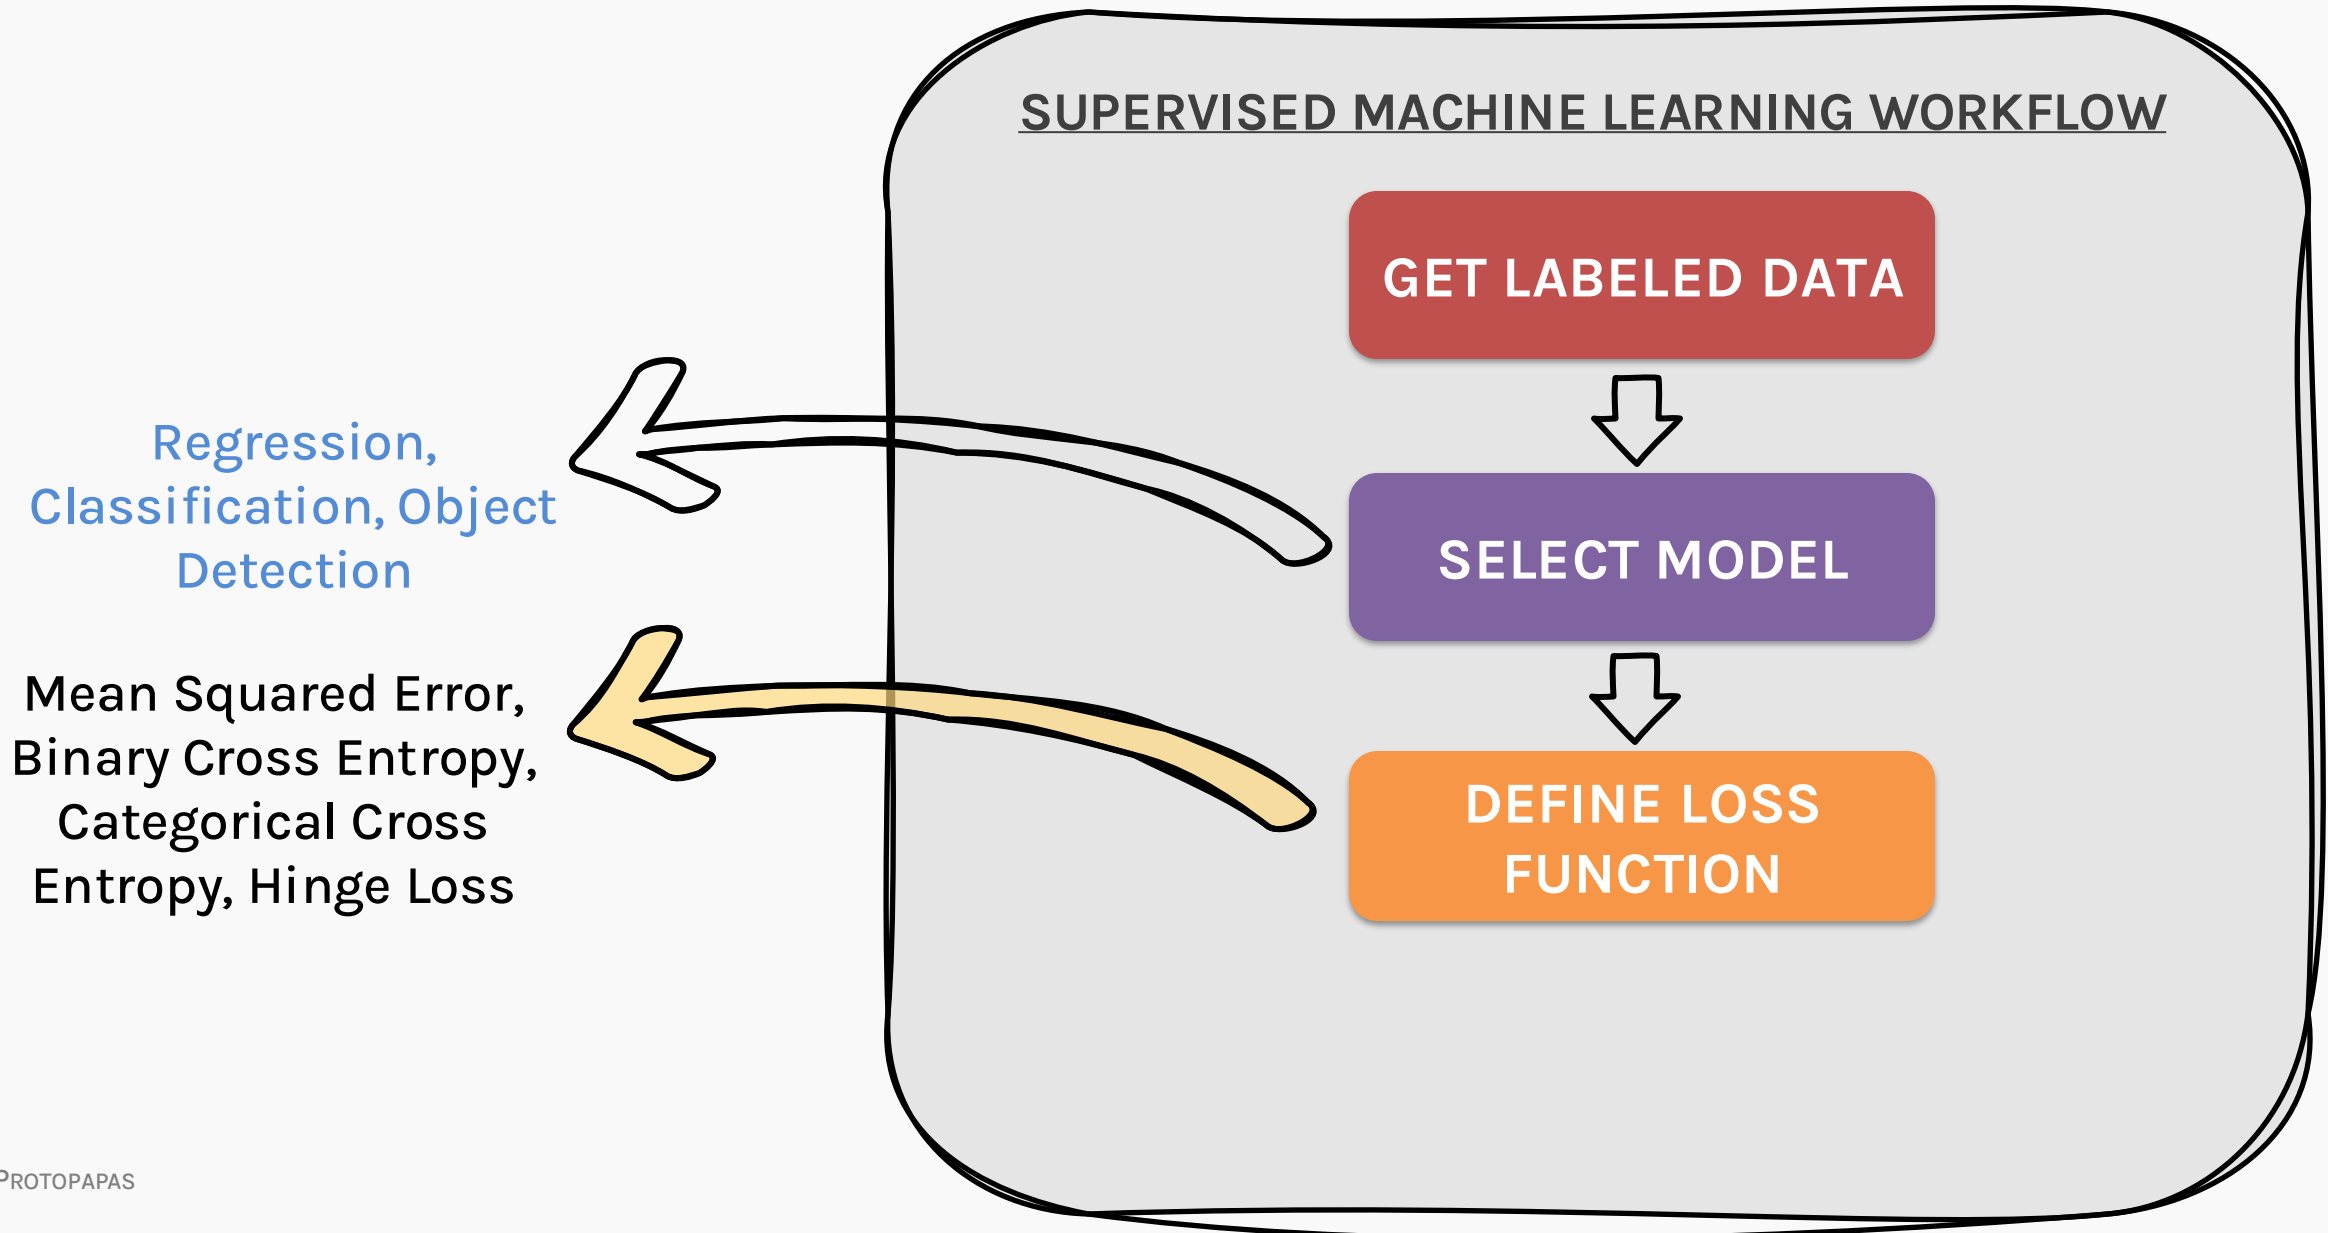
SUPERVISED MACHINE LEARNING WORKFLOW

GET LABELED DATA

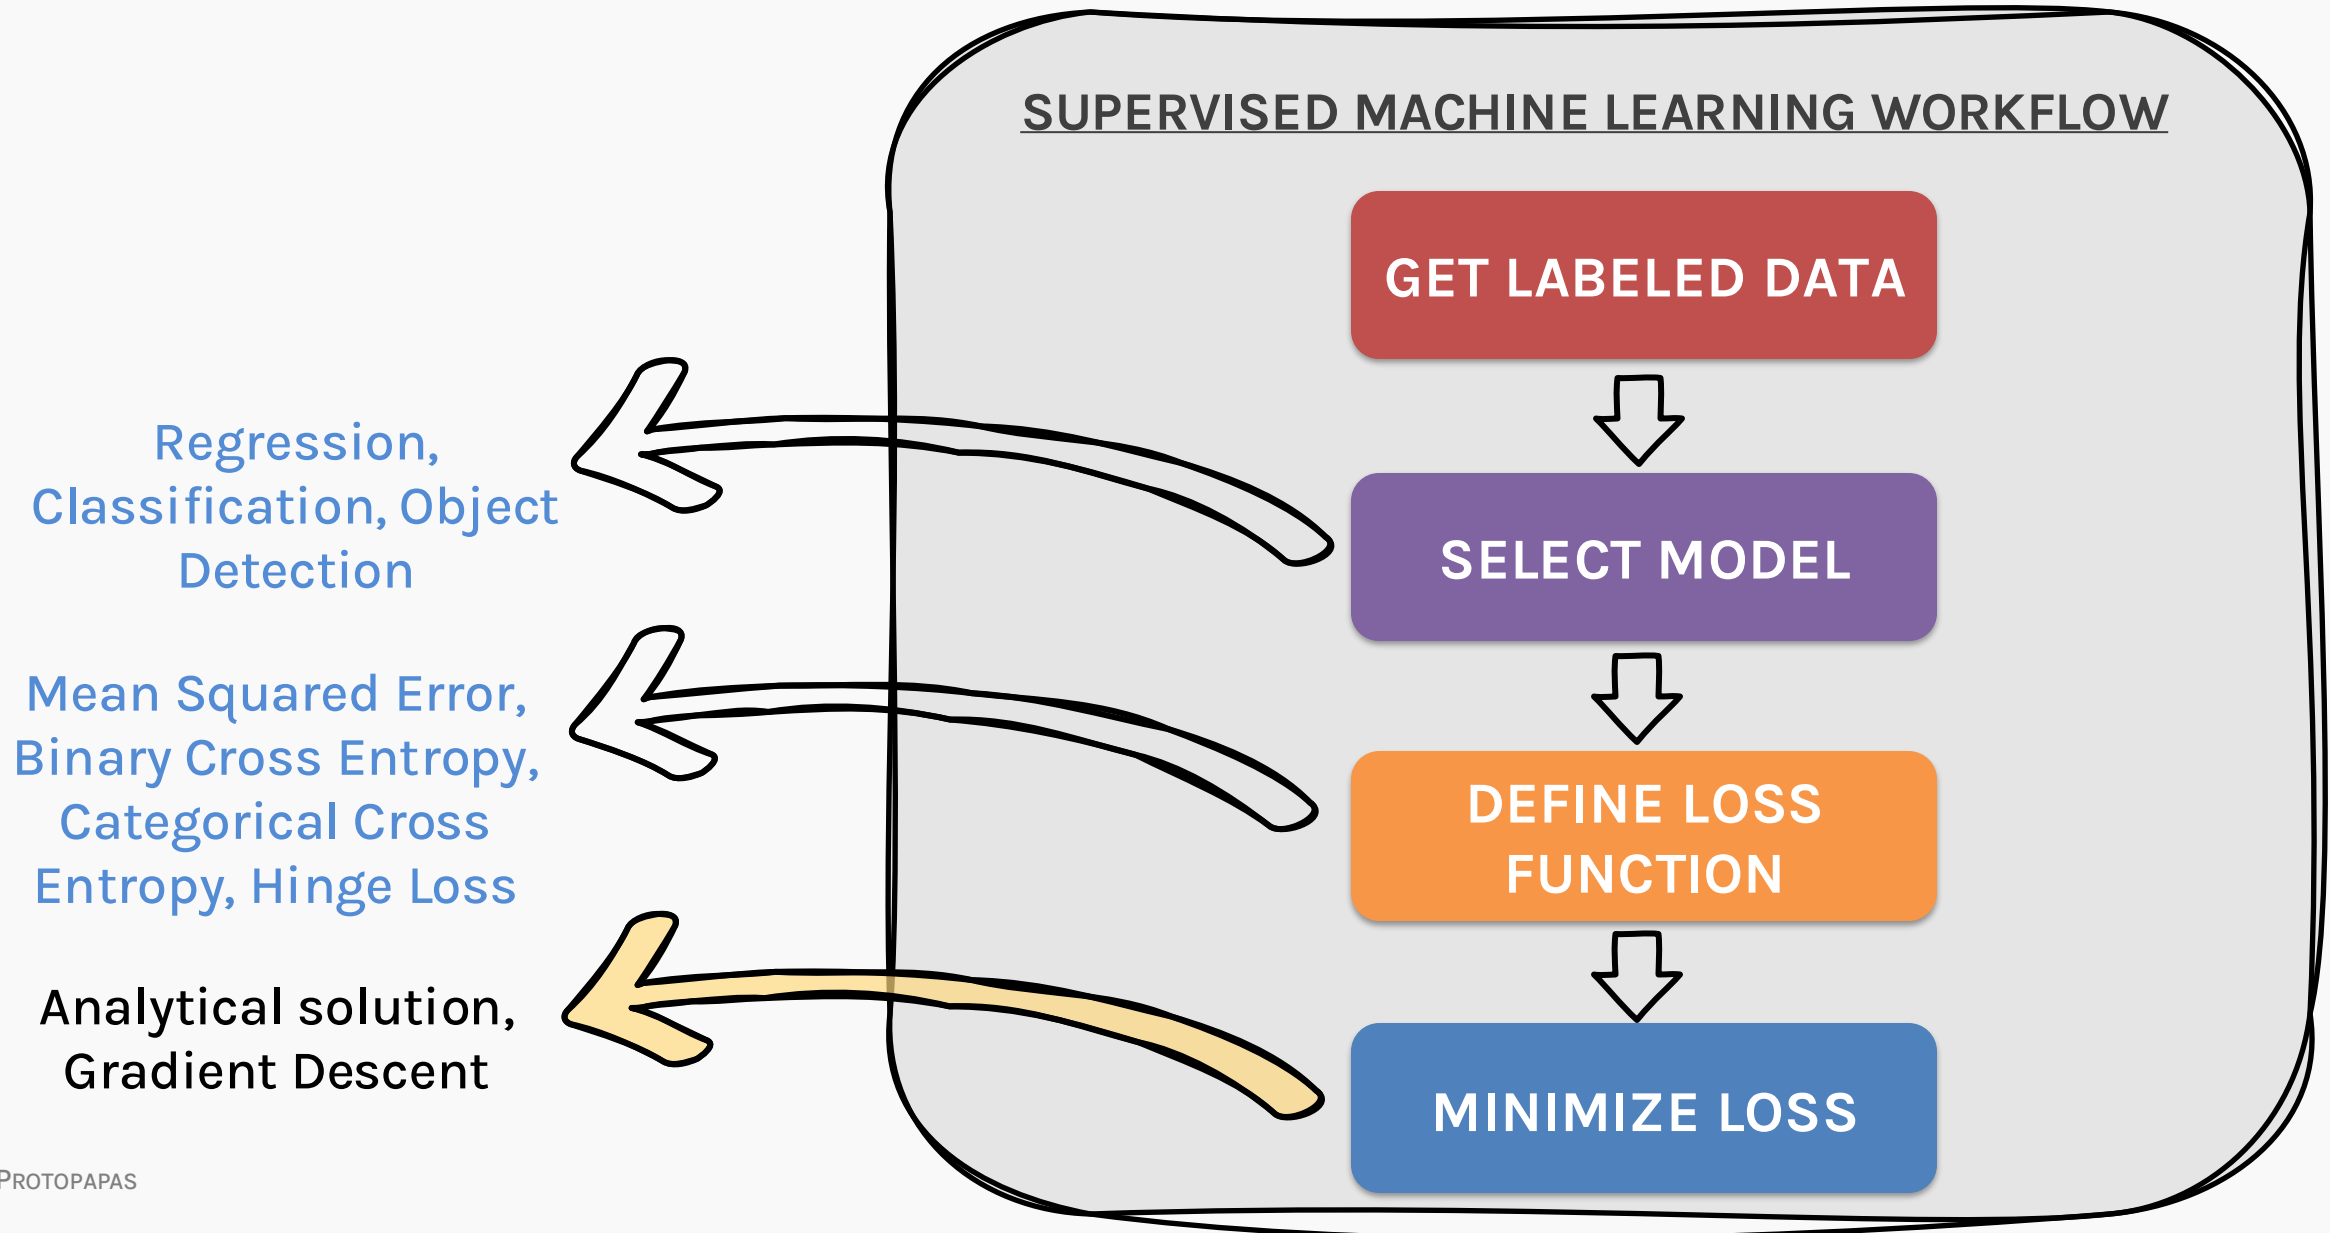
Building blocks of supervised machine learning



Building blocks of supervised machine learning



Building blocks of **supervised** machine learning



Outline

- Introduction
- Review of basic concepts
- **Linear to Logistic Regression**
- Logistic Regression to ANN
- Perceptron - Single neuron network
- Multi-Layer Perceptron (MLP)

Linear to Logistic Regression



Before we dive into what an ANN looks like, let us review [Logistic Regression](#)

Linear to Logistic Regression



Linear Regression is like predicting **how many points** a Basketball Player will score in a game based on past performance!

Logistic Regression, however, predicts whether the same player will score above or below a certain number of points. Basically, a **Yes** or **No** by giving us the probability between **0** and **1**

What is Classification?

Consider the dataset that contains a **binary outcome** AHD for 303 patients who presented with chest pain.

X
predictors

Y
Yes or No
response variable

| Age | Sex | ChestPain | RestBP | Chol | MaxHR | ExAng | Thal | AHD |
|-----|-----|--------------|--------|------|-------|-------|------------|-----|
| 63 | 1 | typical | 145 | 233 | 150 | 0 | fixed | No |
| 67 | 1 | asymptomatic | 160 | 286 | 108 | 1 | normal | Yes |
| 67 | 1 | asymptomatic | 120 | 229 | 129 | 1 | reversable | Yes |
| 37 | 1 | nonanginal | 130 | 250 | 187 | 0 | normal | No |

What is Classification?

Consider the dataset that contains a **binary outcome** AHD for 303 patients who presented with chest pain.

X
predictors

| Age | Sex | ChestPain | RestBP | Chol | MaxHR | ExAng | Thal | AHD |
|-----|-----|--------------|--------|------|-------|-------|------------|-----|
| 63 | 1 | typical | 145 | 233 | 150 | 0 | fixed | No |
| 67 | 1 | asymptomatic | 160 | 286 | 108 | 1 | normal | Yes |
| 67 | 1 | asymptomatic | 120 | 229 | 129 | 1 | reversable | Yes |
| 37 | 1 | nonanginal | 130 | 250 | 187 | 0 | normal | No |

Yes indicates
presence of
heart disease

What is Classification?

Consider the dataset that contains a **binary outcome** AHD for 303 patients who presented with chest pain.

| X predictors | | | | | | | | AHD No indicates absence of heart disease |
|-------------------|-----|--------------|--------|------|-------|-------|------------|---|
| Age | Sex | ChestPain | RestBP | Chol | MaxHR | ExAng | Thal | |
| 63 | 1 | typical | 145 | 233 | 150 | 0 | fixed | |
| 67 | 1 | asymptomatic | 160 | 286 | 108 | 1 | normal | |
| 67 | 1 | asymptomatic | 120 | 229 | 129 | 1 | reversable | |
| 37 | 1 | nonanginal | 130 | 250 | 187 | 0 | normal | No |

Why not Linear Regression?

- Here the response variable y has only two categories.

$$y = \begin{cases} 1, & \text{has heart disease} \\ 0, & \text{no heart disease} \end{cases}$$


- Linear regression could be used to predict the probability $P(y = 1)$ directly from a set of predictors such as sex, cholesterol levels, etc.

If $P(y = 1) \geq 0.5$, we could predict that the patient has heart disease and predict otherwise if $P(y = 1) < 0.5$.

Why not Linear Regression?

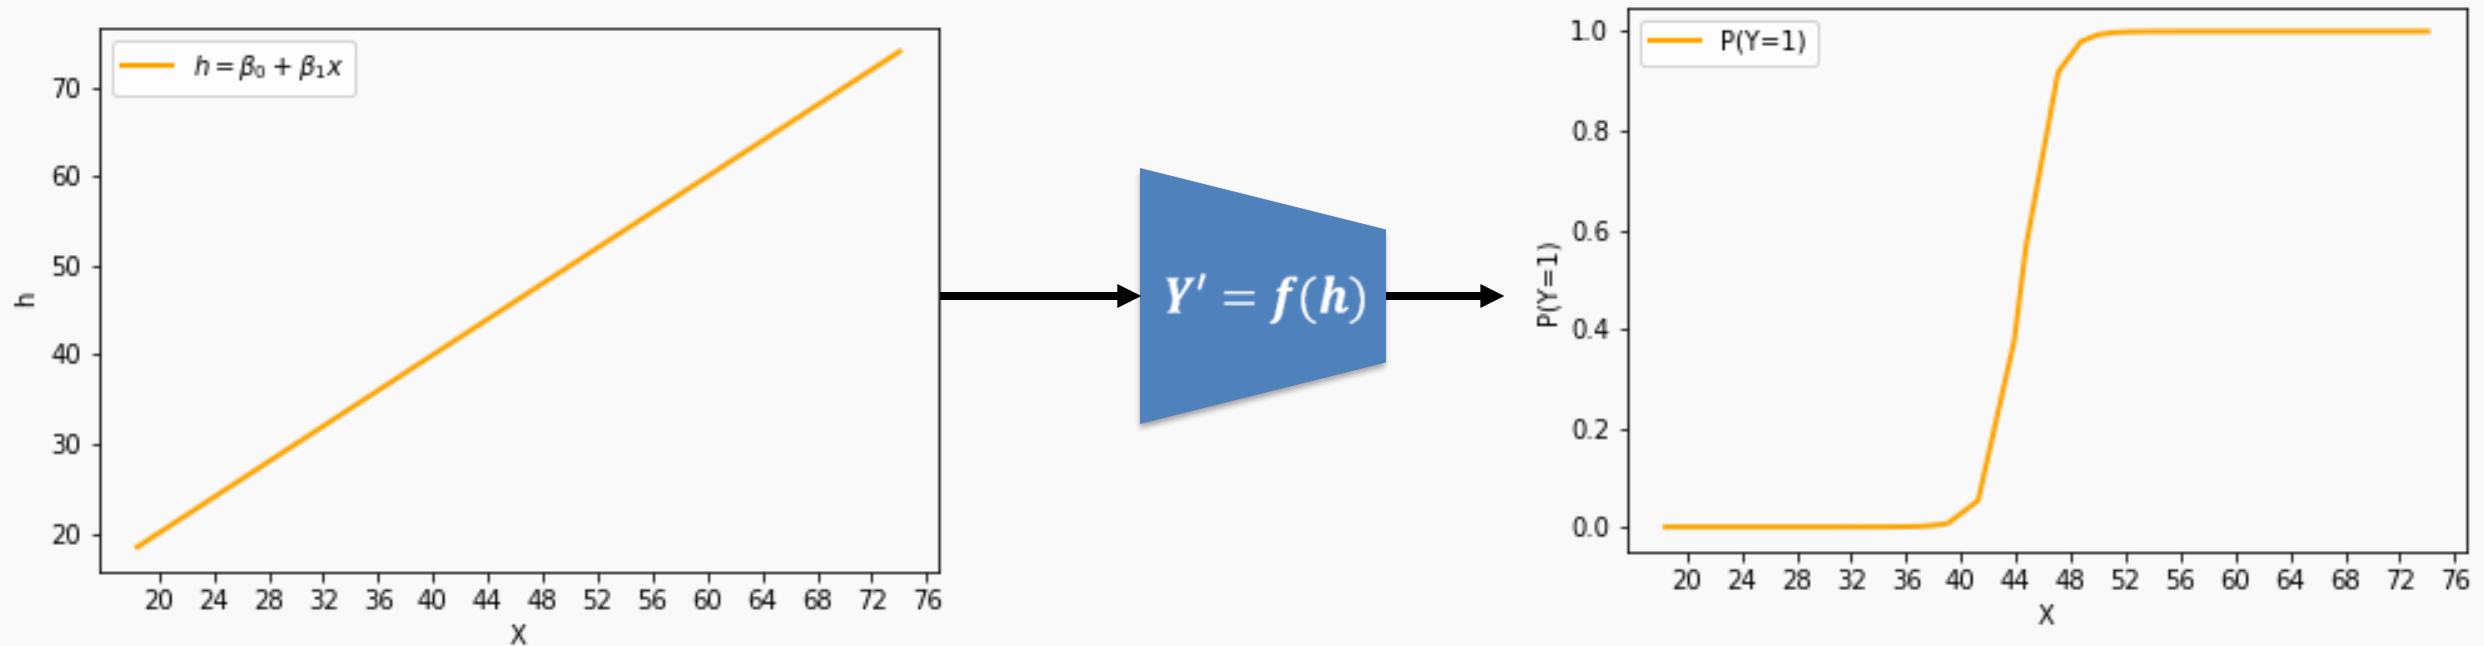
What could go wrong with this linear regression model?

Since this is modeling $P(y = 1)$, values for \hat{y} below 0 and above 1 would not make sense as a probability.




Logistic Regression

Now we know that linear regression yields values for probability that are larger than 1 or smaller than 0. So, what can we do to fix this?



Logistic Regression

We can use the [sigmoid function](#):



Logistic Regression

- Logistic Regression addresses the problem of estimating a probability, $P(y = 1)$, to be outside the range of $[0,1]$.
- The Logistic Regression model uses a function, called the **logistic** function, to model $P(y = 1)$:

$$P(Y = 1) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}} = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X)}}$$

This is the same to the sigmoid function

Using Logistic Regression for Classification

How can we use a logistic regression model to perform classification?

That is, how can we predict when $Y = 1$ vs. when $Y = 0$?

We can classify all observations for which:

- Classify all observations with $\hat{P}(Y = 1) \geq 0.5$ to be in the group associated with $Y = 1$.
- Classify all observations with $\hat{P}(Y = 1) < 0.5$ to be in the group associated with $Y = 0$.

Loss Function

How do we find out if our predictions are any good?

We use a Loss Function!

A Loss Function compares the **predicted output** and the **actual target values**.

In the Linear Regression, we use a
Mean Squared Error

Loss Function

What do we use in Logistic Regression?

We use **Binary Cross Entropy**

The diagram shows the formula for Binary Cross Entropy loss, $L_{BCE} = \sum_i -y_i \log p_i - (1 - y_i) \log(1 - p_i)$, enclosed in a blue rectangular box. Two blue speech bubbles are positioned above the formula: one on the left labeled 'Actual Label' with an arrow pointing to the y_i term, and one on the right labeled 'Predicted Probability' with an arrow pointing to the p_i term. Below the formula box, two arrows point from the $-y_i \log p_i$ and $-(1 - y_i) \log(1 - p_i)$ terms towards a text box at the bottom.

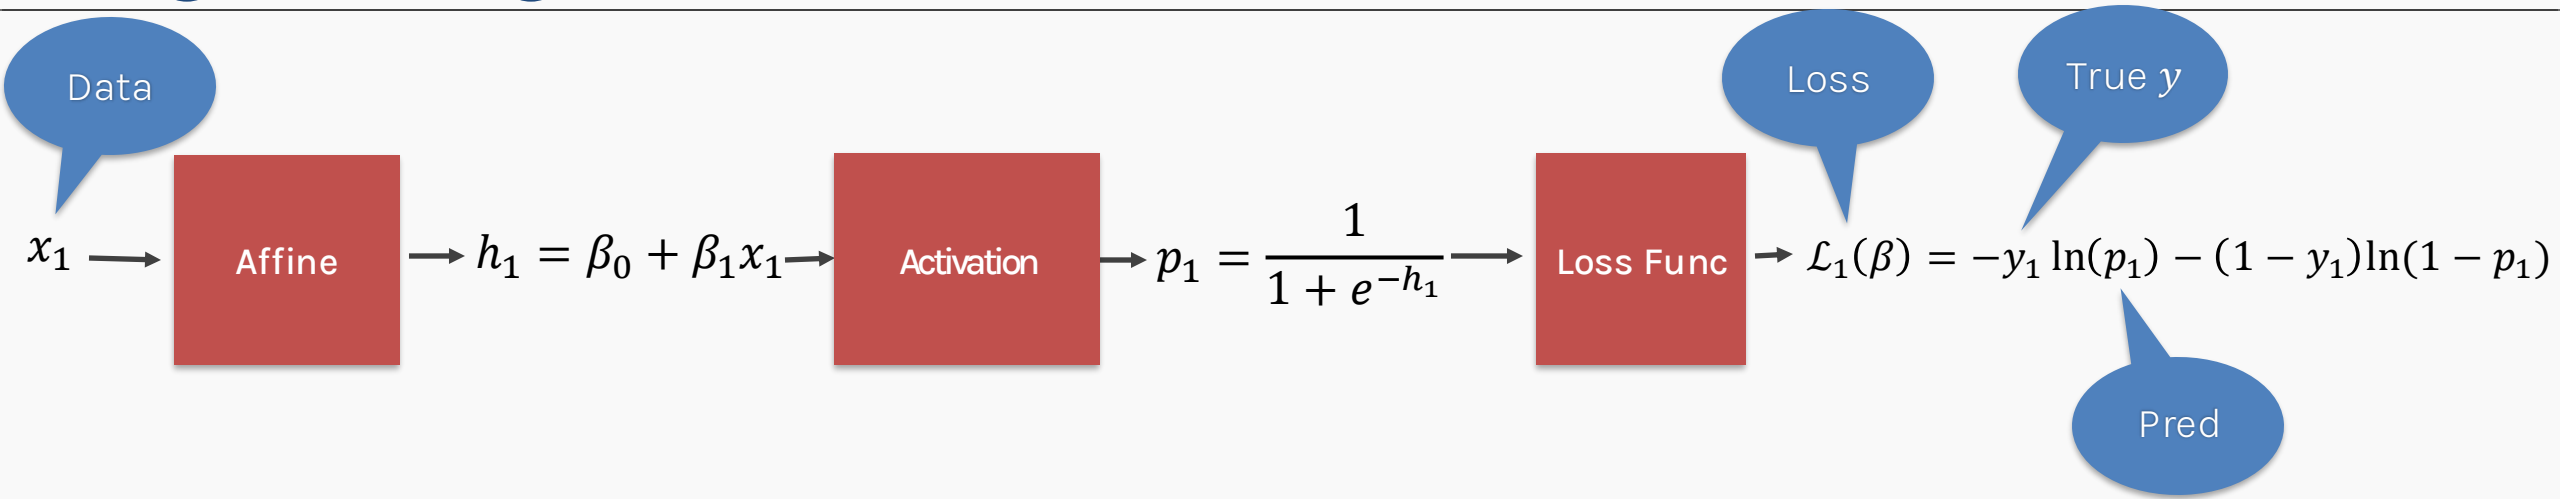
$$L_{BCE} = \sum_i -y_i \log p_i - (1 - y_i) \log(1 - p_i)$$

We subtract, as the goal is to **minimize the loss**


Outline

- Introduction
- Review of basic concepts
- Linear to Logistic Regression
- **Logistic Regression to ANN**
- Perceptron - Single neuron network
- Multi-Layer Perceptron (MLP)

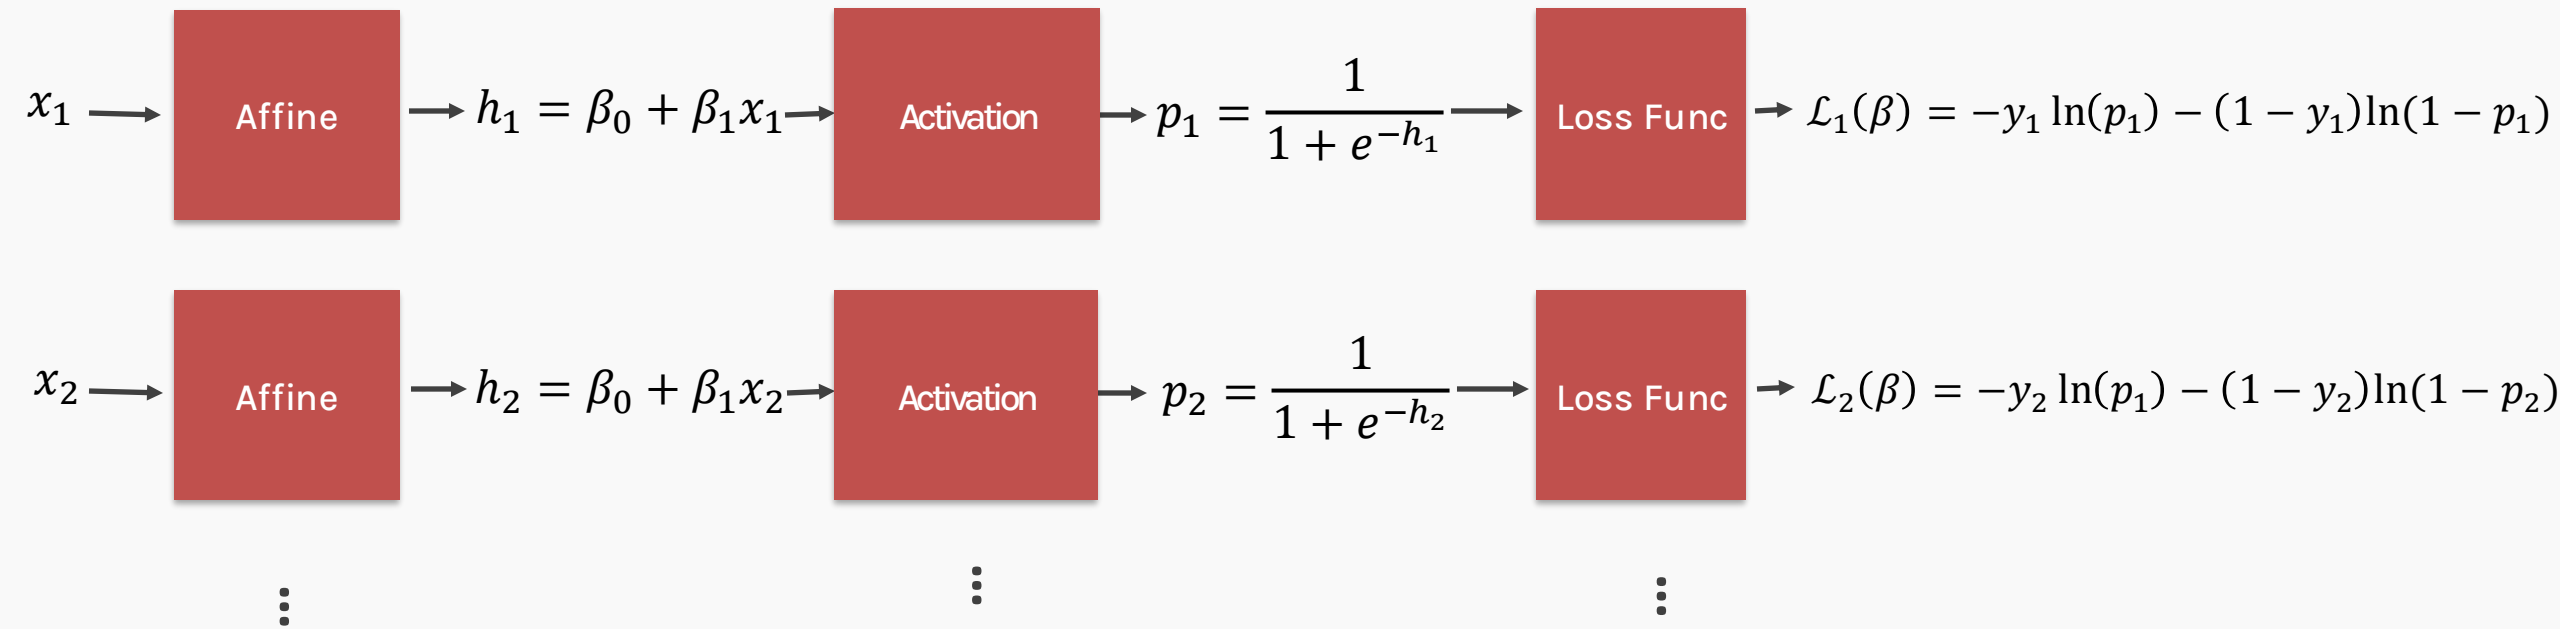
Logistic Regression to ANN




Logistic Regression to ANN



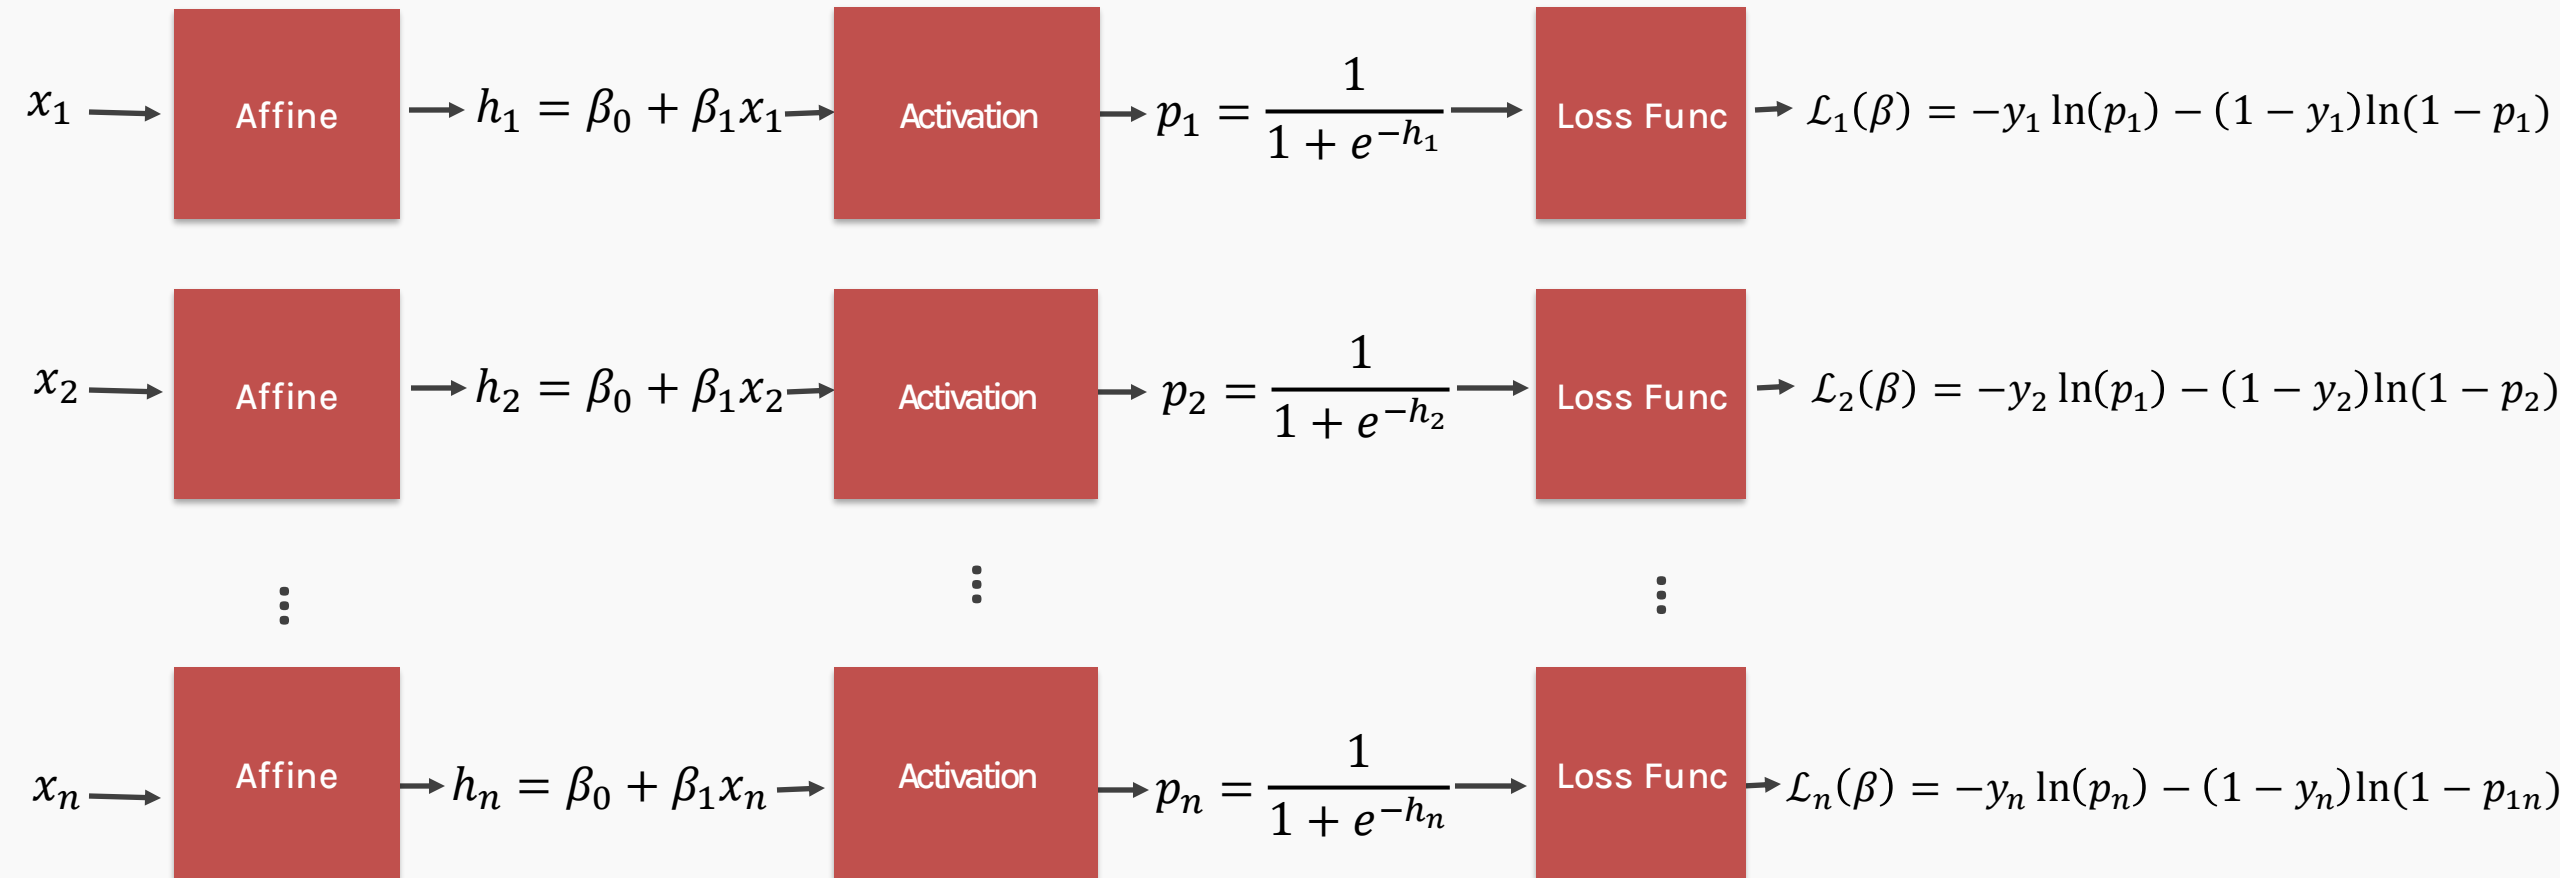
Logistic Regression to ANN



Logistic Regression to ANN



Logistic Regression to ANN



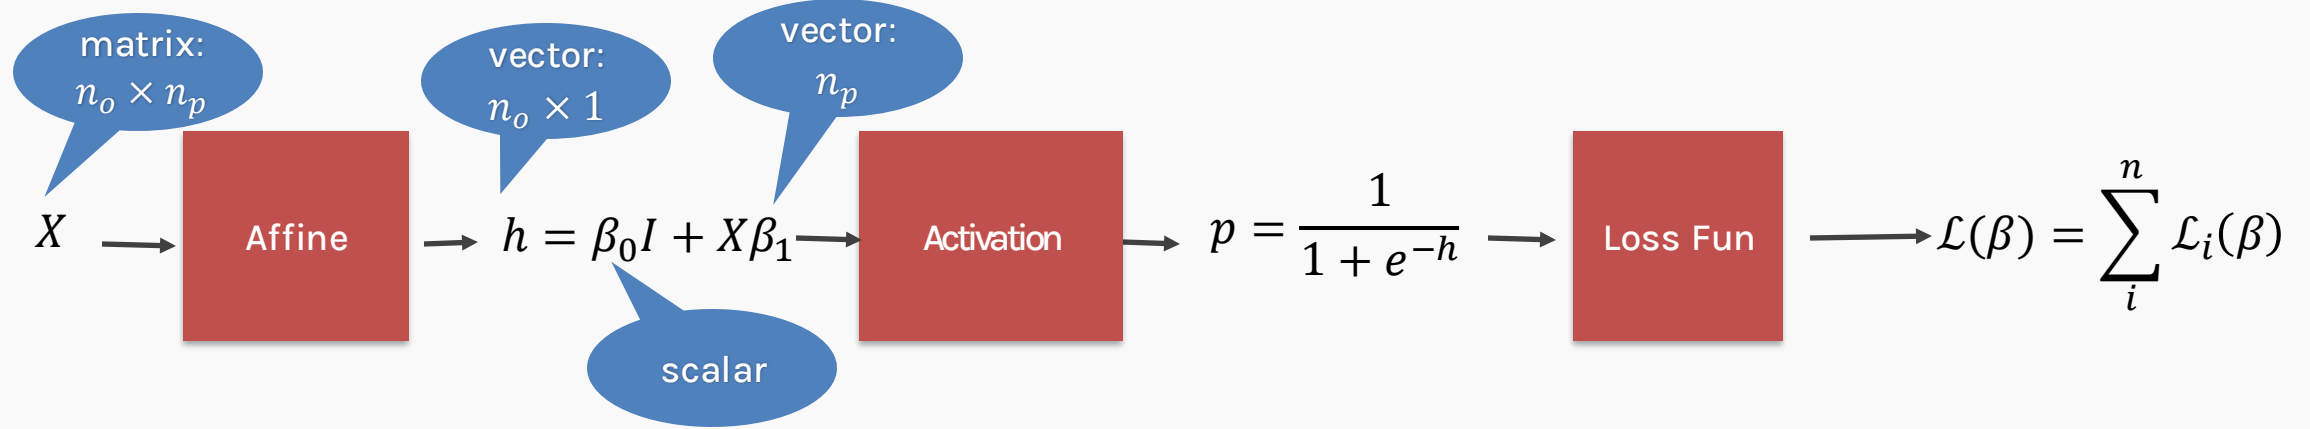
$$\mathcal{L}(\beta) = \sum_i^n \mathcal{L}_i(\beta)$$

Outline

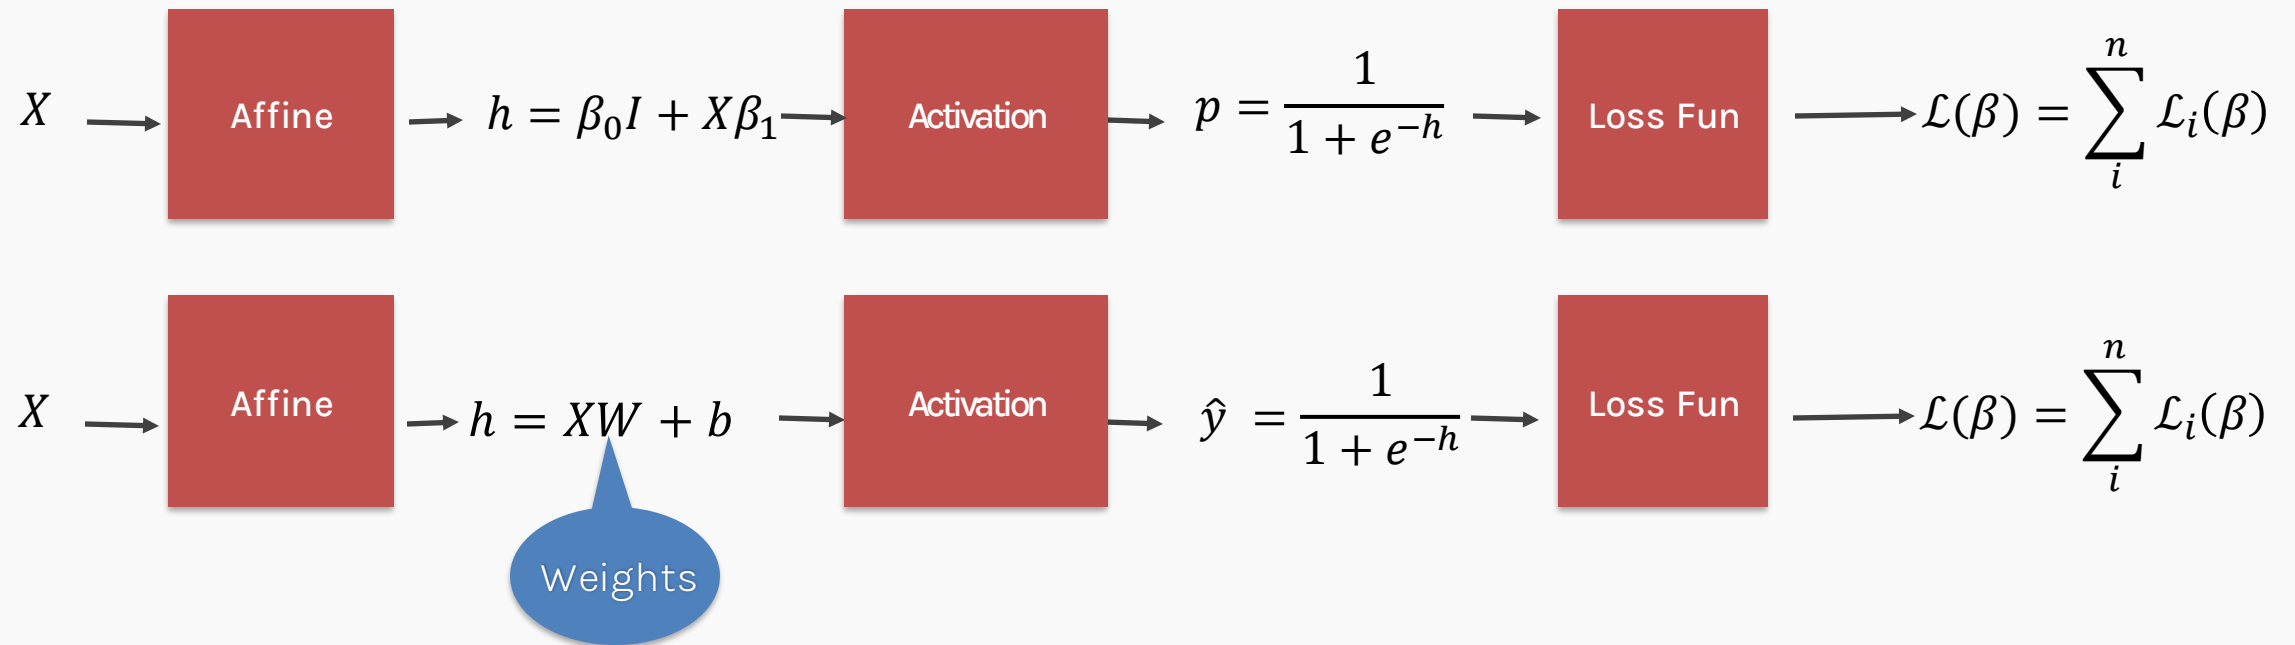
- Introduction
- Review of basic concepts
- Linear to Logistic Regression
- Logistic Regression to ANN
- **Perceptron - Single neuron network**
- Multi-Layer Perceptron (MLP)

Build our first ANN




n_p : number of predictors
 n_o : number of observations



Build our first ANN




Build our first ANN

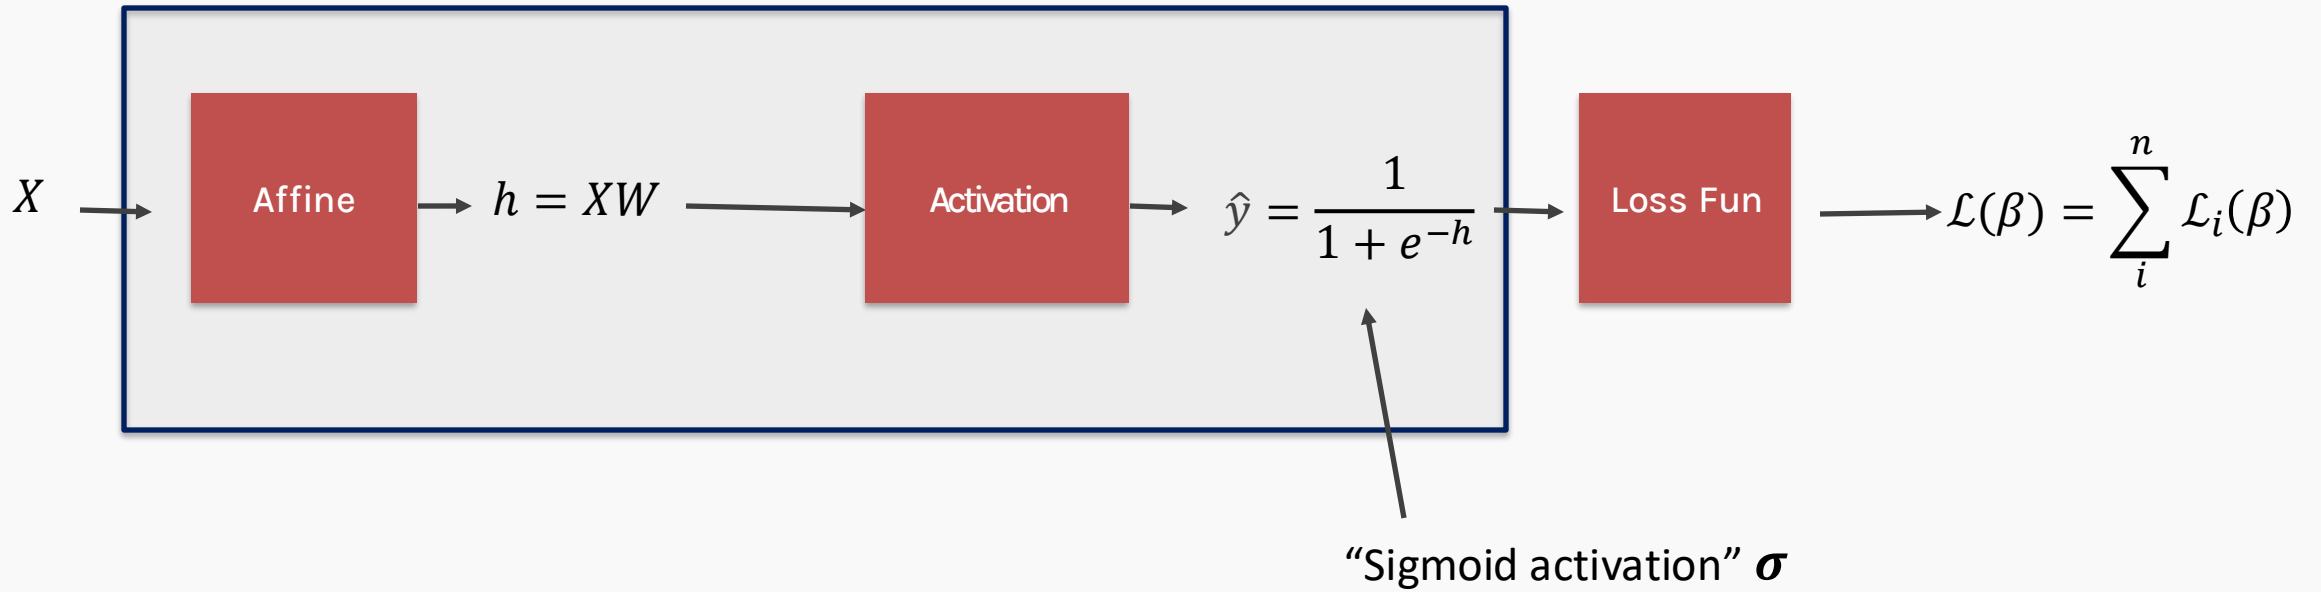


$$X = \begin{bmatrix} 1 & X_{11} & \dots & X_{1p} \\ 1 & \vdots & \dots & \vdots \\ 1 & X_{o1} & \dots & X_{op} \end{bmatrix} \quad W = \begin{bmatrix} b \\ W_1 \\ \vdots \\ W_p \end{bmatrix}$$

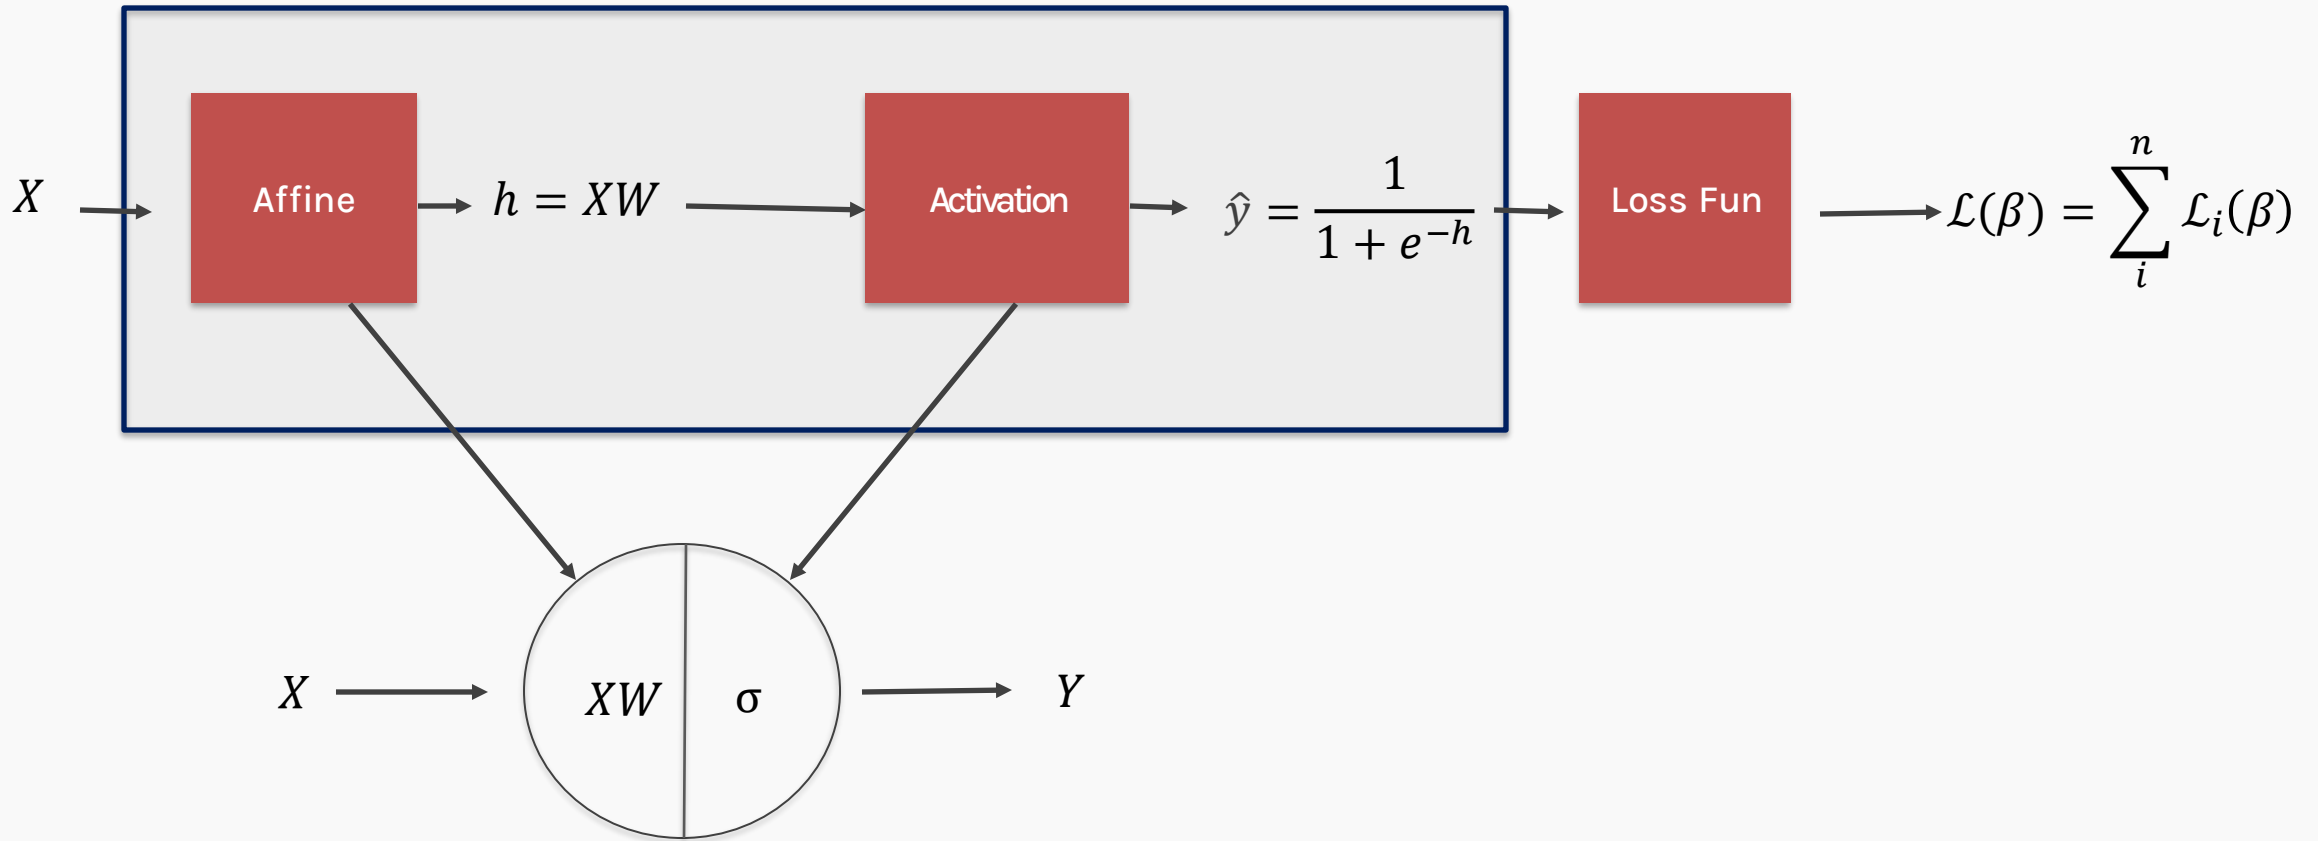
Build our first ANN



Build our first ANN


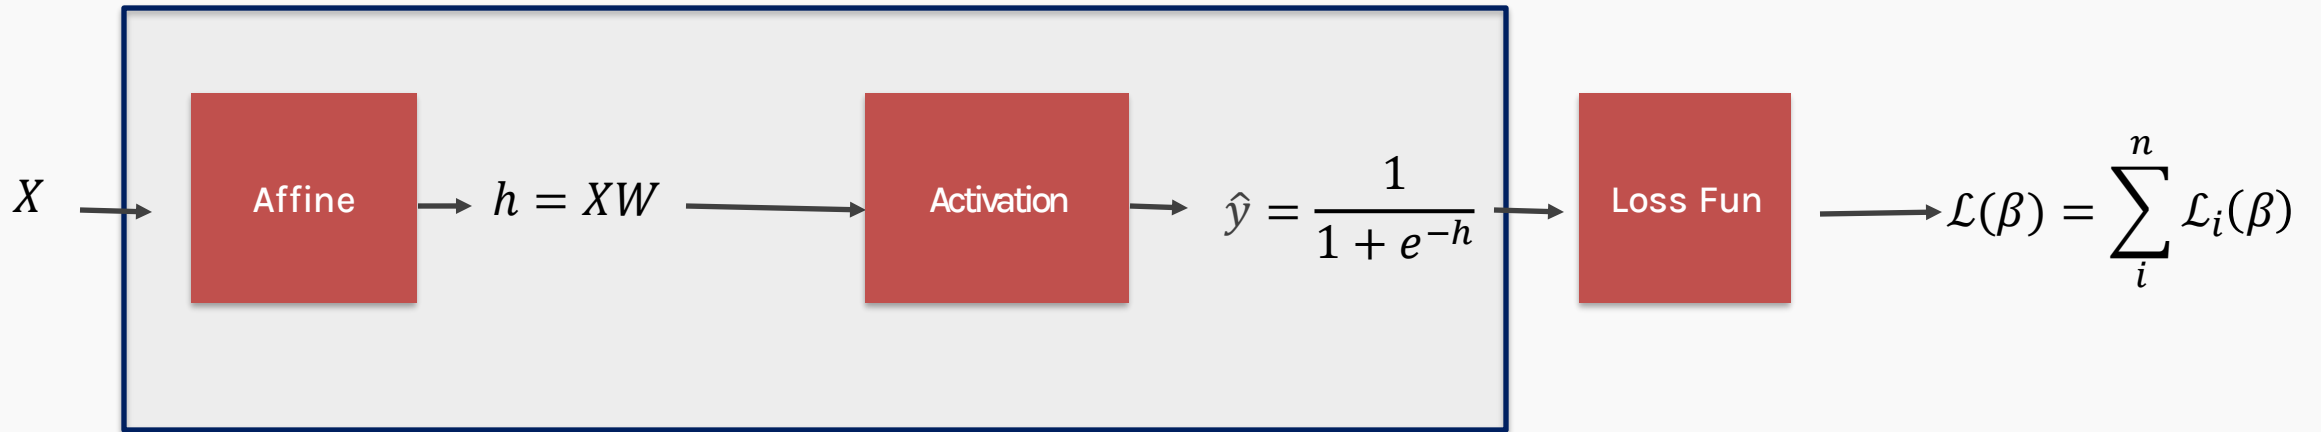


Build our first ANN




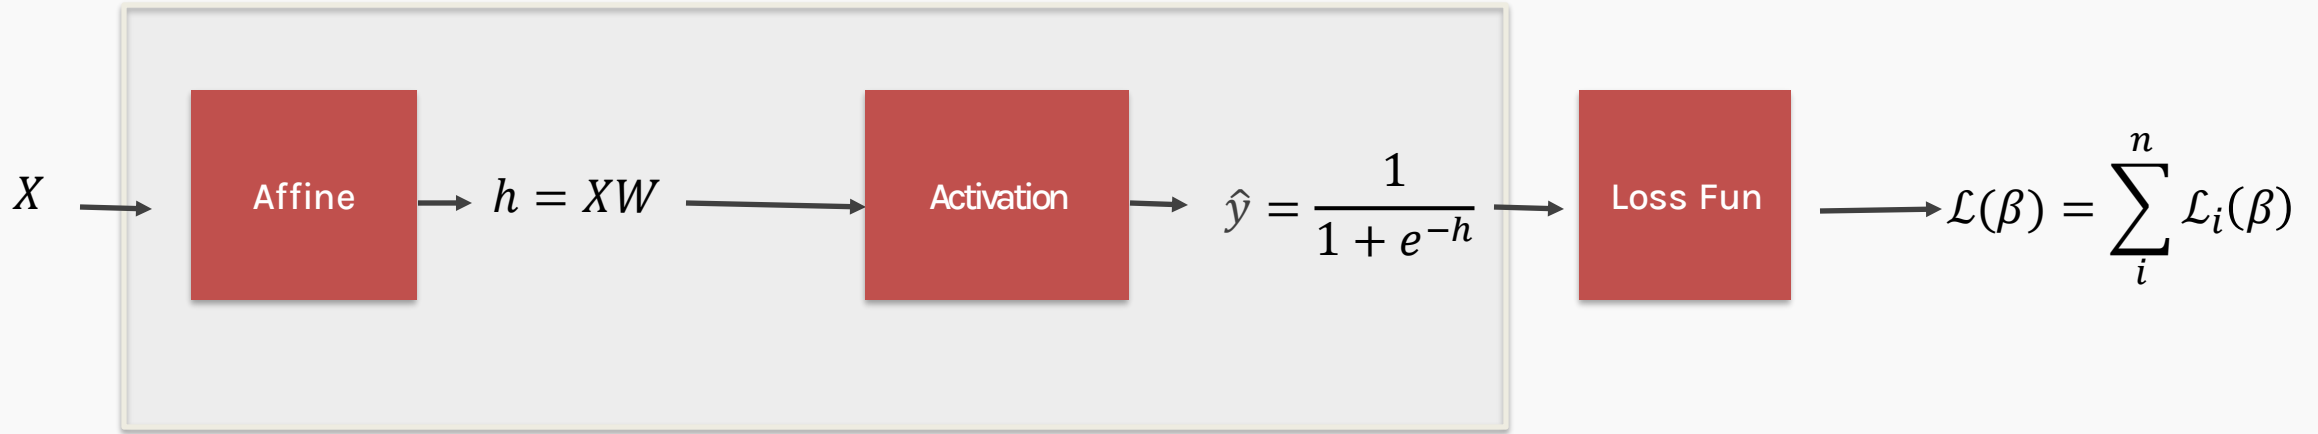
A neuron essentially applies an [affine transformation](#) on the input, followed by an [activation function](#)

Build our first ANN



For simplicity we
can annotate the
weights
outside(W)

Build our first ANN

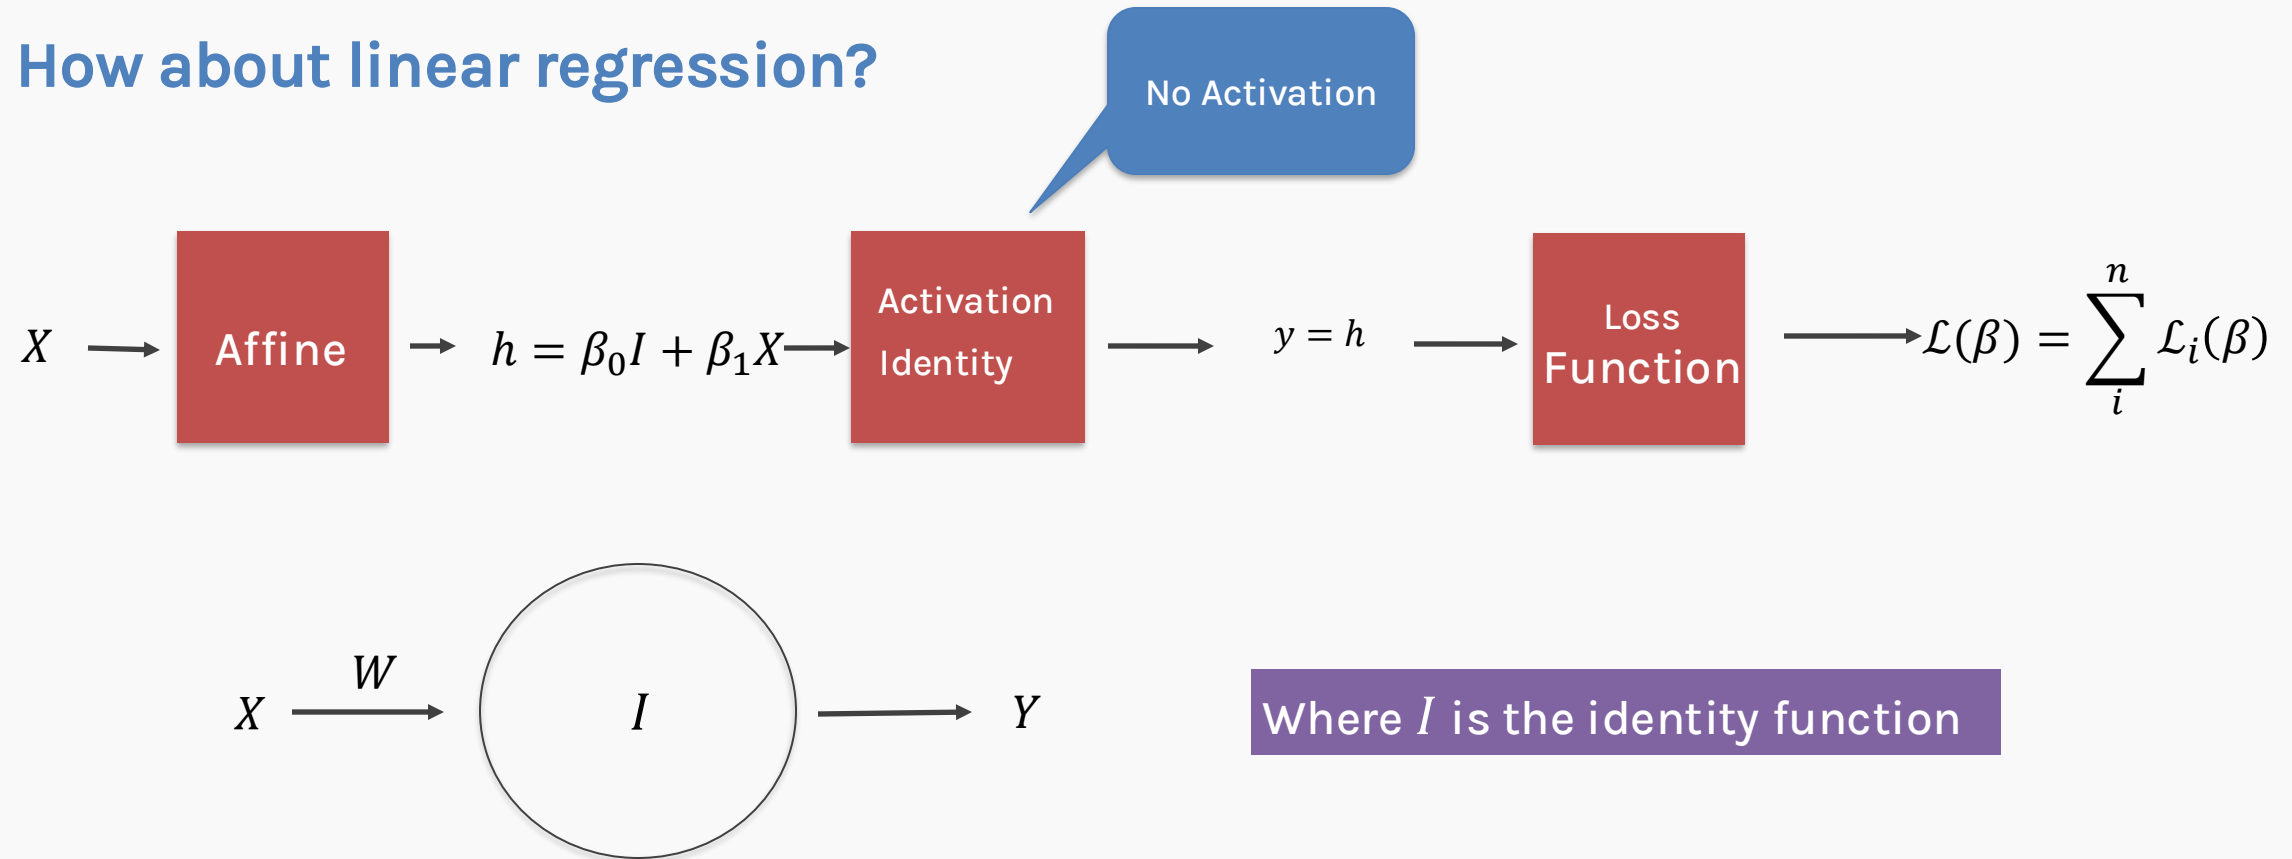


Single Neuron Neural “Network”

A single neuron

Up to this point we just re-branded logistic regression to look like a neuron.

How about linear regression?



Summary

So far:

- A single neuron is simply an **affine** transformation followed by an **activation** function.
- A single neuron with **sigmoid activation** is equivalent to the **logistic regression** and a single neuron with **no activation** is equivalent to **linear regression**.